

```
In [2]: 1 import torch
        2 print("Torch Version: {}".format(torch.__version__))
```

Torch Version: 1.8.1

```
In [3]: 1 import numpy as np
        2 print("Numpy Version: {}".format(np.__version__))
```

Numpy Version: 1.19.2

```
In [4]: 1 import matplotlib
        2 print("Matplotlib Version: {}".format(matplotlib.__version__))
```

Matplotlib Version: 3.3.2

```
In [5]: 1 import torchvision
        2 print("Torchvision Version: {}".format(torchvision.__version__))
```

Torchvision Version: 0.9.1

```
In [6]: 1 import matplotlib.pyplot as plt
        2 import matplotlib.image as mpimg
        3 from torch import nn, optim
        4 import torch.nn.functional as F
        5 from torch.utils.data import DataLoader
        6 from torchvision import datasets, transforms
        7 from torch.utils.data.sampler import SubsetRandomSampler
        8
        9 %matplotlib inline
       10 device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
In [7]: 1 batch_size = 20
        2
        3 valid_size = 0.2
        4
        5 transform = transforms.Compose([
        6     transforms.ToTensor(),
        7     transforms.Normalize((0.5), (0.5,))
        8 ])
        9
       10 train_data = datasets.FashionMNIST('F_MNIST_data/', download=True, train=True, transform=transform)
       11
       12 test_data = datasets.FashionMNIST('F_MNIST_data/', download=True, train=False, transform=transform)
       13
       14 num_train = len(train_data)
       15 indices = list(range(num_train))
       16 np.random.shuffle(indices)
       17 split = int(np.floor(valid_size * num_train))
       18 train_idx, valid_idx = indices[split:], indices[:split]
       19
       20 train_sampler = SubsetRandomSampler(train_idx)
       21 valid_sampler = SubsetRandomSampler(valid_idx)
       22
       23 train_loader = DataLoader(train_data, batch_size=batch_size, sampler=train_sampler)
       24
       25 valid_loader = DataLoader(train_data, batch_size=batch_size, sampler=valid_sampler)
       26
       27 test_loader = DataLoader(test_data, batch_size=batch_size)
       28
       29 classes = ['T-shirt/top',
       30            'Trouser',
       31            'Pullover',
       32            'Dress',
       33            'Coat',
       34            'Sandal',
       35            'Shirt',
       36            'Sneaker',
       37            'Bag',
       38            'Ankle Boot']
```

In [8]: ▶

```
1 class Net(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.fc1 = nn.Linear(784, 512)
5         self.fc2 = nn.Linear(512, 256)
6         self.fc3 = nn.Linear(256, 128)
7         self.fc4 = nn.Linear(128, 64)
8         self.fc5 = nn.Linear(64, 10)
9         self.dropout = nn.Dropout(p=0.2)
10
11     def forward(self, x):
12         x = x.view(x.shape[0], -1)
13
14         x = self.dropout(F.relu(self.fc1(x)))
15         x = self.dropout(F.relu(self.fc2(x)))
16         x = self.dropout(F.relu(self.fc3(x)))
17         x = self.dropout(F.relu(self.fc4(x)))
18         x = F.log_softmax(self.fc5(x), dim=-1)
19
20     return x
```

In [9]: ▶

```
1 model = Net()
2 model = model.to(device)
3 criterion = nn.NLLLoss()
4 optimizer = optim.SGD(model.parameters(), lr=0.01)
```

In [42]:

```

1 epochs = 50
2 train_losses, valid_losses = [], []
3 valid_loss_min = np.Inf
4
5 for e in range(epochs):
6     train_loss, valid_loss = 0.0, 0.0
7     model.train()
8     for images, labels in train_loader:
9         images, labels = images.to(device), labels.to(device)
10
11         optimizer.zero_grad()
12         output = model(images)
13         loss = criterion(output, labels)
14         loss.backward()
15         optimizer.step()
16         train_loss += loss.item()*images.size(0)
17
18     model.eval
19     for images, labels in valid_loader:
20         images, labels = images.to(device), labels.to(device)
21
22         output = model(images)
23         loss = criterion(output, labels)
24         valid_loss += loss.item()*images.size(0)
25
26
27     train_loss = train_loss/len(train_loader.sampler)
28     train_losses.append(train_loss)
29     valid_loss = valid_loss/len(valid_loader.sampler)
30     valid_losses.append(valid_loss)
31
32     print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
33         e+1, train_loss, valid_loss))
34
35     if valid_loss <= valid_loss_min:
36         print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
37             valid_loss_min,
38             valid_loss))
39         torch.save(model.state_dict(), 'fashion_mnist_ann.pt')
40         valid_loss_min = valid_loss

```

```

Epoch: 1      Training Loss: 1.358861      Validation Loss: 0.769927
Validation loss decreased (inf --> 0.769927). Saving model ...
Epoch: 2      Training Loss: 0.663367      Validation Loss: 0.583180
Validation loss decreased (0.769927 --> 0.583180). Saving model ...
Epoch: 3      Training Loss: 0.548222      Validation Loss: 0.521153
Validation loss decreased (0.583180 --> 0.521153). Saving model ...
Epoch: 4      Training Loss: 0.486952      Validation Loss: 0.484718
Validation loss decreased (0.521153 --> 0.484718). Saving model ...
Epoch: 5      Training Loss: 0.452480      Validation Loss: 0.442287
Validation loss decreased (0.484718 --> 0.442287). Saving model ...
Epoch: 6      Training Loss: 0.425699      Validation Loss: 0.414073
Validation loss decreased (0.442287 --> 0.414073). Saving model ...
Epoch: 7      Training Loss: 0.404787      Validation Loss: 0.406415
Validation loss decreased (0.414073 --> 0.406415). Saving model ...
Epoch: 8      Training Loss: 0.388965      Validation Loss: 0.382574
Validation loss decreased (0.406415 --> 0.382574). Saving model ...
Epoch: 9      Training Loss: 0.372344      Validation Loss: 0.371904
Validation loss decreased (0.382574 --> 0.371904). Saving model ...
Epoch: 10     Training Loss: 0.359971      Validation Loss: 0.366394
Validation loss decreased (0.371904 --> 0.366394). Saving model ...
Epoch: 11     Training Loss: 0.348125      Validation Loss: 0.362578
Validation loss decreased (0.366394 --> 0.362578). Saving model ...
Epoch: 12     Training Loss: 0.340430      Validation Loss: 0.354380
Validation loss decreased (0.362578 --> 0.354380). Saving model ...
Epoch: 13     Training Loss: 0.327693      Validation Loss: 0.373747
Epoch: 14     Training Loss: 0.321065      Validation Loss: 0.357246
Epoch: 15     Training Loss: 0.314037      Validation Loss: 0.354296
Validation loss decreased (0.354380 --> 0.354296). Saving model ...
Epoch: 16     Training Loss: 0.304257      Validation Loss: 0.343482
Validation loss decreased (0.354296 --> 0.343482). Saving model ...
Epoch: 17     Training Loss: 0.297523      Validation Loss: 0.338897
Validation loss decreased (0.343482 --> 0.338897). Saving model ...
Epoch: 18     Training Loss: 0.290501      Validation Loss: 0.337632
Validation loss decreased (0.338897 --> 0.337632). Saving model ...
Epoch: 19     Training Loss: 0.285073      Validation Loss: 0.346223
Epoch: 20     Training Loss: 0.279887      Validation Loss: 0.345977
Epoch: 21     Training Loss: 0.273485      Validation Loss: 0.337106
Validation loss decreased (0.337632 --> 0.337106). Saving model ...
Epoch: 22     Training Loss: 0.267907      Validation Loss: 0.339344
Epoch: 23     Training Loss: 0.261199      Validation Loss: 0.347261
Epoch: 24     Training Loss: 0.257007      Validation Loss: 0.349846
Epoch: 25     Training Loss: 0.253143      Validation Loss: 0.323562
Validation loss decreased (0.337106 --> 0.323562). Saving model ...
Epoch: 26     Training Loss: 0.247418      Validation Loss: 0.334285
Epoch: 27     Training Loss: 0.242037      Validation Loss: 0.324759
Epoch: 28     Training Loss: 0.237751      Validation Loss: 0.341492


```

```
Epoch: 29      Training Loss: 0.231133      Validation Loss: 0.321932
Validation loss decreased (0.323562 --> 0.321932). Saving model ...
Epoch: 30      Training Loss: 0.228064      Validation Loss: 0.329132
Epoch: 31      Training Loss: 0.226681      Validation Loss: 0.353209
Epoch: 32      Training Loss: 0.220912      Validation Loss: 0.351970
Epoch: 33      Training Loss: 0.218139      Validation Loss: 0.330462
Epoch: 34      Training Loss: 0.212439      Validation Loss: 0.334524
Epoch: 35      Training Loss: 0.208175      Validation Loss: 0.335599
Epoch: 36      Training Loss: 0.205895      Validation Loss: 0.349774
Epoch: 37      Training Loss: 0.202952      Validation Loss: 0.319005
Validation loss decreased (0.321932 --> 0.319005). Saving model ...
Epoch: 38      Training Loss: 0.198813      Validation Loss: 0.334688
Epoch: 39      Training Loss: 0.196446      Validation Loss: 0.337607
Epoch: 40      Training Loss: 0.192292      Validation Loss: 0.340911
Epoch: 41      Training Loss: 0.191735      Validation Loss: 0.339332
Epoch: 42      Training Loss: 0.185891      Validation Loss: 0.324097
Epoch: 43      Training Loss: 0.183762      Validation Loss: 0.344853
Epoch: 44      Training Loss: 0.181242      Validation Loss: 0.358128
Epoch: 45      Training Loss: 0.178423      Validation Loss: 0.342752
Epoch: 46      Training Loss: 0.173301      Validation Loss: 0.368151
Epoch: 47      Training Loss: 0.171478      Validation Loss: 0.329854
Epoch: 48      Training Loss: 0.167631      Validation Loss: 0.339990
Epoch: 49      Training Loss: 0.166009      Validation Loss: 0.351595
Epoch: 50      Training Loss: 0.166257      Validation Loss: 0.355389
```

```
In [11]: 1 model.load_state_dict(torch.load("fashion_mnist_ann.pt"))
```

```
Out[11]: <All keys matched successfully>
```

```
In [46]: 1 plt.plot(train_losses, label="Train Loss")
2 plt.plot(valid_losses, label="Valid Loss")
3 plt.legend(frameon=False)
4 plt.show()
```



```

In [49]: ▶ 1 test_loss = 0.0
2 class_correct = list(0. for i in range(10))
3 class_total = list(0. for i in range(10))
4
5 model.eval()
6
7 for images, labels in test_loader:
8     images, labels = images.to(device), labels.to(device)
9
10    output = model(images)
11    loss = criterion(output, labels)
12    test_loss += loss.item()*images.size(0)
13    _, pred = torch.max(output, 1)
14    correct = np.squeeze(pred.eq(labels.data.view_as(pred)))
15
16    for i in range(batch_size):
17        label = labels.data[i]
18        class_correct[label] += correct[i].item()
19        class_total[label] += 1
20
21
22 test_loss = test_loss/len(test_loader.dataset)
23 print('Test Loss: {:.6f}\n'.format(test_loss))
24
25 for i in range(10):
26     if class_total[i] > 0:
27         print('Test Accuracy of %5s: %2d%% (%2d/%2d)' % (
28             str(i), 100 * class_correct[i] / class_total[i],
29             np.sum(class_correct[i]), np.sum(class_total[i])))
30     else:
31         print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))
32
33 print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
34     100. * np.sum(class_correct) / np.sum(class_total),
35     np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 0.334308

Test Accuracy of 0: 88% (883/1000)  
 Test Accuracy of 1: 96% (963/1000)  
 Test Accuracy of 2: 83% (830/1000)  
 Test Accuracy of 3: 88% (886/1000)  
 Test Accuracy of 4: 85% (855/1000)  
 Test Accuracy of 5: 96% (968/1000)  
 Test Accuracy of 6: 63% (637/1000)  
 Test Accuracy of 7: 95% (952/1000)  
 Test Accuracy of 8: 97% (975/1000)  
 Test Accuracy of 9: 95% (958/1000)

Test Accuracy (Overall): 89% (8907/10000)

```

In [13]: ▶ 1 images, labels = next(iter(test_loader))
2 images, labels = images.to(device), labels.to(device)
3
4 output = model(images)
5 _, preds = torch.max(output, 1)
6 images = images.cpu().numpy()
7
8 fig = plt.figure(figsize=(25, 4))
9 for idx in range(20):
10     ax = fig.add_subplot(2, 10, idx+1, xticks=[], yticks=[])
11     ax.imshow(np.squeeze(images[idx]), cmap='gray')
12     ax.set_title("{} ({}).format(classes[preds[idx].item()], classes[labels[idx].item()]),
13                 color=("green" if preds[idx]==labels[idx] else "red"))

```

