

```
In [17]: 1 import torch
2 print("Torch Version: {}".format(torch.__version__))
```

Torch Version: 1.8.1

```
In [18]: 1 import numpy as np
2 print("Numpy Version: {}".format(np.__version__))
```

Numpy Version: 1.19.2

```
In [19]: 1 import matplotlib
2 print("Matplotlib Version: {}".format(matplotlib.__version__))
```

Matplotlib Version: 3.3.2

```
In [20]: 1 import torchvision
2 print("Torchvision Version: {}".format(torchvision.__version__))
```

Torchvision Version: 0.9.1

```
In [21]: 1 import matplotlib.pyplot as plt
2 import matplotlib.image as mpimg
3 from torch import nn, optim
4 import torch.nn.functional as F
5 from torch.utils.data import DataLoader
6 from torchvision import datasets, transforms
7 from torch.utils.data.sampler import SubsetRandomSampler
8
9 %matplotlib inline
10 device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
In [22]: 1 batch_size = 20
2
3 valid_size = 0.2
4
5 transform = transforms.Compose([
6     transforms.ToTensor(),
7     transforms.Normalize((0.5), (0.5,))
8 ])
9
10 train_data = datasets.MNIST('MNIST_data', train=True, download=True, transform=transform)
11
12 test_data = datasets.MNIST('MNIST_data', train=False, download=True, transform=transform)
13
14 num_train = len(train_data)
15 indices = list(range(num_train))
16 np.random.shuffle(indices)
17 split = int(np.floor(valid_size * num_train))
18 train_idx, valid_idx = indices[split:], indices[:split]
19
20 train_sampler = SubsetRandomSampler(train_idx)
21 valid_sampler = SubsetRandomSampler(valid_idx)
22
23 train_loader = DataLoader(train_data, batch_size=batch_size, sampler=train_sampler)
24
25 valid_loader = DataLoader(train_data, batch_size=batch_size, sampler=valid_sampler)
26
27 test_loader = DataLoader(test_data, batch_size=batch_size)
```

```
In [23]: 1 class Net(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.fc1 = nn.Linear(784, 256)
5         self.fc2 = nn.Linear(256, 128)
6         self.fc3 = nn.Linear(128, 64)
7         self.fc4 = nn.Linear(64, 10)
8
9     def forward(self, x):
10         x = x.view(x.shape[0], -1) # flattening the input tensor
11
12         x = F.relu(self.fc1(x))
13         x = F.relu(self.fc2(x))
14         x = F.relu(self.fc3(x))
15         x = F.log_softmax(self.fc4(x), dim=1)
16
17         return x
```

In [24]: ▶

```
1 model = Net()  
2 model = model.to(device)  
3 criterion = nn.CrossEntropyLoss()  
4 optimizer = optim.SGD(model.parameters(), lr=0.003)
```

In [25]:

```

1 epochs = 30
2 train_losses, valid_losses = [], []
3 valid_loss_min = np.Inf
4
5 for e in range(epochs):
6     train_loss, valid_loss = 0.0, 0.0
7     model.train()
8     for images, labels in train_loader:
9         images, labels = images.to(device), labels.to(device)
10
11         optimizer.zero_grad()
12         output = model(images)
13         loss = criterion(output, labels)
14         loss.backward()
15         optimizer.step()
16         train_loss += loss.item()*images.size(0)
17
18     model.eval
19     for images, labels in valid_loader:
20         images, labels = images.to(device), labels.to(device)
21
22         output = model(images)
23         loss = criterion(output, labels)
24         valid_loss += loss.item()*images.size(0)
25
26
27     train_loss = train_loss/len(train_loader.sampler)
28     train_losses.append(train_loss)
29     valid_loss = valid_loss/len(valid_loader.sampler)
30     valid_losses.append(valid_loss)
31
32     print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
33         e+1, train_loss, valid_loss))
34
35     if valid_loss <= valid_loss_min:
36         print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
37             valid_loss_min,
38             valid_loss))
39         torch.save(model.state_dict(), 'mnist_ann.pt')
40         valid_loss_min = valid_loss

```

```

Epoch: 1      Training Loss: 1.645108      Validation Loss: 0.756156
Validation loss decreased (inf --> 0.756156). Saving model ...
Epoch: 2      Training Loss: 0.553181      Validation Loss: 0.463151
Validation loss decreased (0.756156 --> 0.463151). Saving model ...
Epoch: 3      Training Loss: 0.388362      Validation Loss: 0.366195
Validation loss decreased (0.463151 --> 0.366195). Saving model ...
Epoch: 4      Training Loss: 0.329024      Validation Loss: 0.319076
Validation loss decreased (0.366195 --> 0.319076). Saving model ...
Epoch: 5      Training Loss: 0.293908      Validation Loss: 0.296290
Validation loss decreased (0.319076 --> 0.296290). Saving model ...
Epoch: 6      Training Loss: 0.267165      Validation Loss: 0.270356
Validation loss decreased (0.296290 --> 0.270356). Saving model ...
Epoch: 7      Training Loss: 0.244221      Validation Loss: 0.249475
Validation loss decreased (0.270356 --> 0.249475). Saving model ...
Epoch: 8      Training Loss: 0.223188      Validation Loss: 0.230348
Validation loss decreased (0.249475 --> 0.230348). Saving model ...
Epoch: 9      Training Loss: 0.204602      Validation Loss: 0.211217
Validation loss decreased (0.230348 --> 0.211217). Saving model ...
Epoch: 10     Training Loss: 0.187142      Validation Loss: 0.197522
Validation loss decreased (0.211217 --> 0.197522). Saving model ...
Epoch: 11     Training Loss: 0.172631      Validation Loss: 0.186119
Validation loss decreased (0.197522 --> 0.186119). Saving model ...
Epoch: 12     Training Loss: 0.160248      Validation Loss: 0.175641
Validation loss decreased (0.186119 --> 0.175641). Saving model ...
Epoch: 13     Training Loss: 0.147675      Validation Loss: 0.173294
Validation loss decreased (0.175641 --> 0.173294). Saving model ...
Epoch: 14     Training Loss: 0.138442      Validation Loss: 0.152355
Validation loss decreased (0.173294 --> 0.152355). Saving model ...
Epoch: 15     Training Loss: 0.129434      Validation Loss: 0.151776
Validation loss decreased (0.152355 --> 0.151776). Saving model ...
Epoch: 16     Training Loss: 0.120464      Validation Loss: 0.142655
Validation loss decreased (0.151776 --> 0.142655). Saving model ...
Epoch: 17     Training Loss: 0.112771      Validation Loss: 0.143752
Epoch: 18     Training Loss: 0.107162      Validation Loss: 0.128059
Validation loss decreased (0.142655 --> 0.128059). Saving model ...
Epoch: 19     Training Loss: 0.100594      Validation Loss: 0.124086
Validation loss decreased (0.128059 --> 0.124086). Saving model ...
Epoch: 20     Training Loss: 0.094637      Validation Loss: 0.121034
Validation loss decreased (0.124086 --> 0.121034). Saving model ...
Epoch: 21     Training Loss: 0.089433      Validation Loss: 0.115384
Validation loss decreased (0.121034 --> 0.115384). Saving model ...
Epoch: 22     Training Loss: 0.084652      Validation Loss: 0.114727
Validation loss decreased (0.115384 --> 0.114727). Saving model ...
Epoch: 23     Training Loss: 0.080095      Validation Loss: 0.110428
Validation loss decreased (0.114727 --> 0.110428). Saving model ...
Epoch: 24     Training Loss: 0.075144      Validation Loss: 0.119690


```

```
Epoch: 25      Training Loss: 0.071518      Validation Loss: 0.105376
Validation loss decreased (0.110428 --> 0.105376). Saving model ...
Epoch: 26      Training Loss: 0.067284      Validation Loss: 0.104210
Validation loss decreased (0.105376 --> 0.104210). Saving model ...
Epoch: 27      Training Loss: 0.064253      Validation Loss: 0.104602
Epoch: 28      Training Loss: 0.060949      Validation Loss: 0.110857
Epoch: 29      Training Loss: 0.058058      Validation Loss: 0.101158
Validation loss decreased (0.104210 --> 0.101158). Saving model ...
Epoch: 30      Training Loss: 0.054616      Validation Loss: 0.100849
Validation loss decreased (0.101158 --> 0.100849). Saving model ...
```

```
In [26]: 1 model.load_state_dict(torch.load('mnist_ann.pt', map_location=torch.device(device)))
```

Out[26]: <All keys matched successfully>

```
In [30]: 1 plt.plot(train_losses, label="Train Loss")
2 plt.plot(valid_losses, label="Valid Loss")
3 plt.legend(frameon=False)
4 plt.show()
```



```

In [28]: 1 test_loss = 0.0
2 class_correct = list(0. for i in range(10))
3 class_total = list(0. for i in range(10))
4
5 model.eval()
6
7 for images, labels in test_loader:
8     images, labels = images.to(device), labels.to(device)
9
10    output = model(images)
11    loss = criterion(output, labels)
12    test_loss += loss.item()*images.size(0)
13    _, pred = torch.max(output, 1)
14    correct = np.squeeze(pred.eq(labels.data.view_as(pred)))
15
16    for i in range(batch_size):
17        label = labels.data[i]
18        class_correct[label] += correct[i].item()
19        class_total[label] += 1
20
21
22 test_loss = test_loss/len(test_loader.dataset)
23 print('Test Loss: {:.6f}\n'.format(test_loss))
24
25 for i in range(10):
26     if class_total[i] > 0:
27         print('Test Accuracy of %2s: %2d%% (%2d/%2d)' % (
28             str([i]), 100 * class_correct[i] / class_total[i],
29             np.sum(class_correct[i]), np.sum(class_total[i])))
30     else:
31         print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))
32
33 print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
34     100. * np.sum(class_correct) / np.sum(class_total),
35     np.sum(class_correct), np.sum(class_total)))

```

Test Loss: 0.090483

Test Accuracy of [0]: 98% (965/980)  
 Test Accuracy of [1]: 99% (1127/1135)  
 Test Accuracy of [2]: 97% (1011/1032)  
 Test Accuracy of [3]: 95% (964/1010)  
 Test Accuracy of [4]: 97% (957/982)  
 Test Accuracy of [5]: 96% (858/892)  
 Test Accuracy of [6]: 98% (942/958)  
 Test Accuracy of [7]: 96% (991/1028)  
 Test Accuracy of [8]: 97% (945/974)  
 Test Accuracy of [9]: 95% (960/1009)

Test Accuracy (Overall): 97% (9720/10000)

```
In [34]: 1 images, labels = next(iter(train_loader))
2 image, labels = images.to(device), labels.to(device)
3
4 img = images[1].view(1,784)
5
6 with torch.no_grad():
7     log_ps = model(img)
8
9 ps = F.softmax(log_ps, dim=1)
10 fig = plt.figure(figsize=(6,9))
11 ax1 = plt.subplot(121)
12 ax1.imshow(img.resize_(1, 28, 28).cpu().numpy().squeeze())
13 ax1.axis('off')
14
15 ax2 = plt.subplot(122)
16 ax2.set_yticks(np.arange(10))
17 ax2.set_xticks(np.arange(0,1.1,0.25))
18 ax2.set_xlim(0,1.1)
19 ax2.set_aspect(0.1)
20 ax2.set_title('Class Probability')
21 plt.barh(np.arange(10),ps.data.cpu().numpy().squeeze())
22
23 plt.tight_layout()
24
```

