

```
In [23]: 1 import torch
2 print("Torch Version: {}".format(torch.__version__))
```

Torch Version: 1.8.1

```
In [25]: 1 import numpy as np
2 print("Numpy Version: {}".format(np.__version__))
```

Numpy Version: 1.19.2

```
In [27]: 1 import matplotlib
2 print("Matplotlib Version: {}".format(matplotlib.__version__))
```

Matplotlib Version: 3.3.2

```
In [29]: 1 import torchvision
2 print("Torchvision Version: {}".format(torchvision.__version__))
```

Torchvision Version: 0.9.1

```
In [31]: 1 import matplotlib.pyplot as plt
2 import matplotlib.image as mpimg
3 from torch import nn, optim
4 import torch.nn.functional as F
5 from torch.utils.data import DataLoader
6 from torchvision import datasets, transforms
7 from torch.utils.data.sampler import SubsetRandomSampler
8
9 %matplotlib inline
10 device = 'cuda' if torch.cuda.is_available() else 'cpu'
```

```
In [33]: 1 batch_size = 20
2
3 valid_size = 0.2
4
5 transform = transforms.Compose([
6     transforms.ToTensor(),
7     transforms.Normalize((0.5), (0.5,))
8 ])
9
10 train_data = datasets.MNIST('MNIST_data', train=True, download=True, transform=transform)
11
12 test_data = datasets.MNIST('MNIST_data', train=False, download=True, transform=transform)
13
14 num_train = len(train_data)
15 indices = list(range(num_train))
16 np.random.shuffle(indices)
17 split = int(np.floor(valid_size * num_train))
18 train_idx, valid_idx = indices[split:], indices[:split]
19
20 train_sampler = SubsetRandomSampler(train_idx)
21 valid_sampler = SubsetRandomSampler(valid_idx)
22
23 train_loader = DataLoader(train_data, batch_size=batch_size, sampler=train_sampler)
24
25 valid_loader = DataLoader(train_data, batch_size=batch_size, sampler=valid_sampler)
26
27 test_loader = DataLoader(test_data, batch_size=batch_size)
```

```
In [ ]: 1 class Net(nn.Module):
2     def __init__(self):
3         super().__init__()
4         self.fc1 = nn.Linear(784, 256)
5         self.fc2 = nn.Linear(256, 128)
6         self.fc3 = nn.Linear(128, 64)
7         self.fc4 = nn.Linear(64, 10)
8
9     def forward(self, x):
10         x = x.view(x.shape[0], -1) # flattening the input tensor
11
12         x = F.relu(self.fc1(x))
13         x = F.relu(self.fc2(x))
14         x = F.relu(self.fc3(x))
15         x = F.log_softmax(self.fc4(x), dim=1)
16
17         return x
```

```
In [32]: ▶ 1 model = Classifier()  
2 model = model.to(device)  
3 criterion = nn.CrossEntropyLoss()  
4 optimizer = optim.SGD(model.parameters(), lr=0.003)
```

In [36]:

```

1 epochs = 30
2 train_losses, valid_losses = [], []
3 valid_loss_min = np.Inf
4
5 for e in range(epochs):
6     train_loss, valid_loss = 0.0, 0.0
7     model.train()
8     for images, labels in train_loader:
9         images, labels = images.to(device), labels.to(device)
10
11         optimizer.zero_grad()
12         output = model(images)
13         loss = criterion(output, labels)
14         loss.backward()
15         optimizer.step()
16         train_loss += loss.item()*images.size(0)
17
18     model.eval
19     for images, labels in valid_loader:
20         images, labels = images.to(device), labels.to(device)
21
22         output = model(images)
23         loss = criterion(output, labels)
24         valid_loss += loss.item()*images.size(0)
25
26
27     train_loss = train_loss/len(train_loader.sampler)
28     train_losses.append(train_loss)
29     valid_loss = valid_loss/len(valid_loader.sampler)
30     valid_losses.append(valid_loss)
31
32     print('Epoch: {} \tTraining Loss: {:.6f} \tValidation Loss: {:.6f}'.format(
33         e+1, train_loss, valid_loss))
34
35     if valid_loss <= valid_loss_min:
36         print('Validation loss decreased ({:.6f} --> {:.6f}). Saving model ...'.format(
37             valid_loss_min,
38             valid_loss))
39         torch.save(model.state_dict(), 'mnist_ann.pt')
40         valid_loss_min = valid_loss

```

```

Epoch: 1      Training Loss: 0.331936      Validation Loss: 0.316636
Validation loss decreased (inf --> 0.316636). Saving model ...
Epoch: 2      Training Loss: 0.303024      Validation Loss: 0.296594
Validation loss decreased (0.316636 --> 0.296594). Saving model ...
Epoch: 3      Training Loss: 0.278292      Validation Loss: 0.267787
Validation loss decreased (0.296594 --> 0.267787). Saving model ...
Epoch: 4      Training Loss: 0.255892      Validation Loss: 0.249430
Validation loss decreased (0.267787 --> 0.249430). Saving model ...
Epoch: 5      Training Loss: 0.234609      Validation Loss: 0.236686
Validation loss decreased (0.249430 --> 0.236686). Saving model ...
Epoch: 6      Training Loss: 0.214459      Validation Loss: 0.216773
Validation loss decreased (0.236686 --> 0.216773). Saving model ...
Epoch: 7      Training Loss: 0.196832      Validation Loss: 0.212171
Validation loss decreased (0.216773 --> 0.212171). Saving model ...
Epoch: 8      Training Loss: 0.180280      Validation Loss: 0.181335
Validation loss decreased (0.212171 --> 0.181335). Saving model ...
Epoch: 9      Training Loss: 0.165568      Validation Loss: 0.177877
Validation loss decreased (0.181335 --> 0.177877). Saving model ...
Epoch: 10     Training Loss: 0.152678      Validation Loss: 0.161196
Validation loss decreased (0.177877 --> 0.161196). Saving model ...
Epoch: 11     Training Loss: 0.141372      Validation Loss: 0.144571
Validation loss decreased (0.161196 --> 0.144571). Saving model ...
Epoch: 12     Training Loss: 0.130932      Validation Loss: 0.144412
Validation loss decreased (0.144571 --> 0.144412). Saving model ...
Epoch: 13     Training Loss: 0.122724      Validation Loss: 0.128777
Validation loss decreased (0.144412 --> 0.128777). Saving model ...
Epoch: 14     Training Loss: 0.114138      Validation Loss: 0.124194
Validation loss decreased (0.128777 --> 0.124194). Saving model ...
Epoch: 15     Training Loss: 0.107077      Validation Loss: 0.120858
Validation loss decreased (0.124194 --> 0.120858). Saving model ...
Epoch: 16     Training Loss: 0.100580      Validation Loss: 0.112536
Validation loss decreased (0.120858 --> 0.112536). Saving model ...
Epoch: 17     Training Loss: 0.094357      Validation Loss: 0.112285
Validation loss decreased (0.112536 --> 0.112285). Saving model ...
Epoch: 18     Training Loss: 0.088760      Validation Loss: 0.105653
Validation loss decreased (0.112285 --> 0.105653). Saving model ...
Epoch: 19     Training Loss: 0.084454      Validation Loss: 0.103083
Validation loss decreased (0.105653 --> 0.103083). Saving model ...
Epoch: 20     Training Loss: 0.079355      Validation Loss: 0.099246
Validation loss decreased (0.103083 --> 0.099246). Saving model ...
Epoch: 21     Training Loss: 0.074687      Validation Loss: 0.094391
Validation loss decreased (0.099246 --> 0.094391). Saving model ...
Epoch: 22     Training Loss: 0.070907      Validation Loss: 0.095798
Epoch: 23     Training Loss: 0.067106      Validation Loss: 0.091535
Validation loss decreased (0.094391 --> 0.091535). Saving model ...
Epoch: 24     Training Loss: 0.063286      Validation Loss: 0.094253

```

```


Epoch: 25      Training Loss: 0.060219      Validation Loss: 0.092203
Epoch: 26      Training Loss: 0.057201      Validation Loss: 0.092112
Epoch: 27      Training Loss: 0.053932      Validation Loss: 0.092661
Epoch: 28      Training Loss: 0.052097      Validation Loss: 0.086928
Validation loss decreased (0.091535 --> 0.086928). Saving model ...
Epoch: 29      Training Loss: 0.048747      Validation Loss: 0.083981
Validation loss decreased (0.086928 --> 0.083981). Saving model ...
Epoch: 30      Training Loss: 0.046583      Validation Loss: 0.084904

```

```
In [37]: 1 model.load_state_dict(torch.load('mnist_ann.pt', map_location=torch.device(device)))
```

Out[37]: <All keys matched successfully>

```
In [39]: 1 plt.plot(train_losses, label='Training loss')
2 plt.plot(valid_losses, label='Test loss')
3 plt.legend(frameon=False)
4 plt.show()
```



```
In [56]: 1 test_loss = 0.0
2 class_correct = list(0. for i in range(10))
3 class_total = list(0. for i in range(10))
4
5 model.eval()
6
7 for images, labels in test_loader:
8     images, labels = images.to(device), labels.to(device)
9
10    output = model(images)
11    loss = criterion(output, labels)
12    test_loss += loss.item()*images.size(0)
13    _, pred = torch.max(output, 1)
14    correct = np.squeeze(pred.eq(labels.data.view_as(pred)))
15
16    for i in range(batch_size):
17        label = labels.data[i]
18        class_correct[label] += correct[i].item()
19        class_total[label] += 1
20
21
22 test_loss = test_loss/len(test_loader.dataset)
23 print('Test Loss: {:.6f}\n'.format(test_loss))
24
25 for i in range(10):
26     if class_total[i] > 0:
27         print('Test Accuracy of %2s: %2d%% (%2d/%2d)' % (
28             str([i]), 100 * class_correct[i] / class_total[i],
29             np.sum(class_correct[i]), np.sum(class_total[i])))
30     else:
31         print('Test Accuracy of %5s: N/A (no training examples)' % (classes[i]))
32
33 print('\nTest Accuracy (Overall): %2d%% (%2d/%2d)' % (
34     100. * np.sum(class_correct) / np.sum(class_total),
35     np.sum(class_correct), np.sum(class_total)))
```

Test Loss: 0.082571

```

Test Accuracy of [0]: 98% (966/980)
Test Accuracy of [1]: 99% (1125/1135)
Test Accuracy of [2]: 97% (1006/1032)
Test Accuracy of [3]: 97% (987/1010)
Test Accuracy of [4]: 96% (952/982)
Test Accuracy of [5]: 97% (871/892)
Test Accuracy of [6]: 97% (936/958)
Test Accuracy of [7]: 95% (980/1028)
Test Accuracy of [8]: 97% (948/974)
Test Accuracy of [9]: 96% (978/1009)

```

Test Accuracy (Overall): 97% (9749/10000)

```
In [77]: 1 images, labels = next(iter(train_loader))
2 image, labels = images.to(device), labels.to(device)
3
4 img = images[1].view(1,784)
5
6 with torch.no_grad():
7     log_ps = model(img)
8
9 ps = F.softmax(log_ps, dim=1)
10 fig = plt.figure(figsize=(6,9))
11 ax1 = plt.subplot(121)
12 ax1.imshow(img.resize_(1, 28, 28).cpu().numpy().squeeze())
13 ax1.axis('off')
14
15 ax2 = plt.subplot(122)
16 ax2.set_yticks(np.arange(10))
17 ax2.set_xticks(np.arange(0,1.1,0.25))
18 ax2.set_xlim(0,1.1)
19 ax2.set_aspect(0.1)
20 ax2.set_title('Class Probability')
21 plt.barh(np.arange(10),ps.data.cpu().numpy().squeeze())
22
23 plt.tight_layout()
24
```

